

EECE 5830 Project Phase 6 (TCP) Design File

Introduction

Transmission control protocol (TCP) consists of a number of key features. These include connection setup, connection teardown, dynamic window sizing, dynamic timeouts, internet checksums, flow control, and congestion control.

Please note that while I try to show as much code as possible here, the actual implementation of TCP is much more verbose and *tcpnet.py* should likely be viewed as reference in parallel with this document.

File Descriptions

`tcpnet.py`

Contains the code which defines the protocol implementation.

`unittests.py`

Extensive tests ranging from typical to edge cases to ensure the robustness of the protocol implementation.

`tcptest.py`

Called by the graphical user interface (GUI) to send and receive a file numerous times.

`gui.py`

Contains the code defining the interactable GUI.

Feature Explanation

All of the key TCP features are fully implemented.

TCP Header

```
251     def make_hdr(self, seq_num: int, ack_num: int, checksum: int, flags: int = 0b00000000):
252         hdr_len = 0 # Header length = Header length field value x 4 bytes
253         urg_ptr = 0
254
255         header: bytearray = bytearray(self.SOURCE_PORT.to_bytes(2, 'big') + self.DEST_PORT.to_bytes(2, 'big') +
256                                     seq_num.to_bytes(4, 'big') + ack_num.to_bytes(4, 'big') + hdr_len.to_bytes(1, 'big') + flags.to_bytes(1,
257                                     'big') + self.rx_win_size.to_bytes(2, 'big') + checksum.to_bytes(2, 'big') + urg_ptr.to_bytes(2, 'big') +
258                                     time.time_ns().to_bytes(8, 'big'))
259
260         return header
```

Figure 1

The header follows a standard TCP header format. 8 bytes of the options field has been used to store timestamp information.

Connection Setup

```
131 def _handshake(self, flags = 0):
132     if self.handshake_complete:
133         return
134
135     if flags == 0b0 and not self.handshake_begun:
136         self._handshake_syn()
137     elif flags == 0b000010: # SYN
138         self.handshake_begun = True
139         # print(self.whois, 'Got SYN.', self.last_rxed_seq_num, self.last_rxed_ack_num)
140         self._handshake_syn_ack()
141     elif flags == 0b010010: # SYN-ACK
142         self.handshake_begun = True
143         # print(self.whois, 'Got SYN-ACK.', self.last_rxed_seq_num, self.last_rxed_ack_num)
144         self._handshake_ack()
145     elif flags == 0b010000:
146         # print(self.whois, 'Got ACK.', self.last_rxed_seq_num, self.last_rxed_ack_num)
147         if self.sent_syn_ack:
148             self.handshake_complete = True
149             # print(self.whois, 'HANDSHAKE COMPLETED', self.last_rxed_seq_num, self.last_rxed_ack_num)
150
151     # CLIENT STEP 1 (Part 1)
152     def _handshake_syn(self):
153         # print(self.whois, 'Sending SYN')
154         self.sent_syn = True
155         self.curr_seq_num = 1
156         self.curr_ack_num = 0
157         syn_pkt: bytearray = bytearray(self.make_hdr(seq_num=self.curr_seq_num, ack_num=self.curr_ack_num,
158             flags=0b000010, checksum=0)) # 2
159         self._udt_send(syn_pkt)
160
161     # SERVER STEP 1 (Part 2)
162     def _handshake_syn_ack(self):
163         # print(self.whois, 'Sending SYN-ACK')
164         self.sent_syn_ack = True
165         self.curr_seq_num = 2
166         self.curr_ack_num = self.last_rxed_seq_num + 1
167         self.zero_index = self.curr_ack_num
168         syn_ack_pkt: bytearray = bytearray(self.make_hdr(seq_num=self.curr_seq_num, ack_num=self.curr_ack_num,
169             flags=0b010010, checksum=0)) # 18
170         self._udt_send(syn_ack_pkt)
171
172     # CLIENT STEP 2 (Part 3)
173     def _handshake_ack(self):
174         # print(self.whois, 'Sending ACK')
175         self.sent_ack = True
176         self.curr_seq_num = self.last_rxed_ack_num
177         self.curr_ack_num = self.last_rxed_seq_num + 1
178         self.zero_index = self.curr_seq_num
179         ack_pkt: bytearray = bytearray(self.make_hdr(seq_num=self.curr_seq_num, ack_num=self.curr_ack_num,
180             flags=0b010000, checksum=0)) # 16
181         self.handshake_complete = True
182         # print(self.whois, 'HANDSHAKE COMPLETE', self.curr_seq_num, self.curr_ack_num)
183         self._udt_send(ack_pkt)
```

Figure 2

Connection setup is conducted via the TCP three-way handshake, facilitated by the functions `_handshake`, `_handshake_syn`, `_handshake_syn_ack`, and `_handshake_ack` (see: Figure 2). When the TCPNet class object is instantiated, the `_tcp_rx_thread` thread, which handles all incoming packets, is started. If it receives a handshake SYN packet, it responds with a SYN-ACK packet and sets its sequence and ack numbers appropriately (see: Figure 3).

```

464         if not self.handshake_complete: # Handshake is incomplete.
465             if self.send_data is not None: # We are the sender.
466                 if flags is None: # We listened but heard nothing and are the sender.
467                     self._handshake(0) # Fire the initial handshake.
468                 else:
469                     self._handshake(flags)
470             elif flags is not None and flags > 0: # We are the receiver (or not ready to send) and got a flag.
471                 self._handshake(flags)

```

Figure 3

The original SYN would have been sent if a send data request was made. In our example, the sender now receives a SYN-ACK and responds with an ACK. Once this handshake is completed by both sides, the data transfer can begin.

Note: Handshake failure is unlikely or impossible since on a timeout, the last sent packet is re-sent.

Connection Teardown

```

182     def _teardown(self):
183         self._teardown_fin()
184
185     def _teardown_fin(self):
186         # 4-way handshake
187         # print(self.whois, 'TEARDOWN HAS BEEN CALLED!')
188         self.teardown_initiated = True
189
190         fin_pkt: bytearray = bytearray(self.make_hdr(seq_num=self.curr_seq_num, ack_num=self.curr_ack_num,
191                                                     flags=0b000001, checksum=0))
192         self._udt_send(fin_pkt)
193
194         # self.done = True
195
196     def _teardown_ack(self):
197         # print(self.whois, 'TEARDOWN HAS BEEN REQUESTED!')
198
199         ack_pkt: bytearray = bytearray(self.make_hdr(seq_num=self.curr_seq_num, ack_num=self.curr_ack_num,
200                                                     flags=0b010000, checksum=0))
201
202         self._udt_send(ack_pkt)
203
204         if not self.teardown_initiated:
205             # If teardown initiated, then we must be the initiator and have already sent a FIN.
206             fin_pkt: bytearray = bytearray(self.make_hdr(seq_num=self.curr_seq_num, ack_num=self.curr_ack_num,
207                                                         flags=0b000001, checksum=0))
208             self._udt_send(fin_pkt)
209
210         self.teardown_initiated = True
211         # self._shutdown()
212         self.done = True
213         # self.log_final()
214         # print('self.done?', self.done)

```

Figure 4

Connection teardown is handled via TCP's four-way handshake and is facilitated by the functions `_teardown`, `_teardown_fin`, and `_teardown_ack` (see: Figure 4). A teardown is initiated by the sender when the last byte it wants to send has been acknowledged. It then sends a FIN

packet, which is received and responded to with an ACK and another FIN packet. At this point both ends of the connection close (but remain instantiated for another connection if desired).

Dynamic Window Sizing and Congestion Control

```
305     def _handle_winsize(self, timeout):
306         ack = self.last_rxed_ack_num
307         seq = self.curr_seq_num
308
309         # This is where we reset win_size to 1/2 due to packet loss.
310         if ack != seq + self.MAX_DATA_SIZE:
311             self.consecutive_nacks += 1
312         else:
313             self.consecutive_nacks = 0
314         if timeout:
315             self.consecutive_nacks = 0
316             self.rx_win_size = 1
317         elif self.consecutive_nacks > 3:
318             self.consecutive_nacks = 0
319             self.rx_win_size = int(self.rx_win_size / 2)
320             if self.rx_win_size < 1:
321                 self.rx_win_size = 1
322         else:
323             self.rx_win_size += 1
```

Figure 5

Dynamic window sizing is handled by `_handle_winsize`, which is called by the receiving thread each timeout or packet receive (see: Figure 5). This implementation follows the TCP Tahoe protocol, where multiple consecutive NACKs result in a halving of the window size, and timeouts result in resetting the window size to 1. Otherwise, the window size is increased.

Dynamic Timeouts and Flow Control

```
405     # Calculates the updated timeout based on RTT.
406     if timestamp is not None:
407         sample_rtt = (time.time_ns() - timestamp) / 1e9
408         self.estimated_rtt = (1 - TCPNet.TYPICAL_RTT) * self.estimated_rtt + TCPNet.TYPICAL_RTT * sample_rtt
409         self.dev_rtt = (1 - TCPNet.TYPICAL_BETA) * self.dev_rtt + TCPNet.TYPICAL_BETA * abs(sample_rtt - self.estimated_rtt)
410         self.timeout_interval = self.estimated_rtt + 4 * self.dev_rtt
411         # print('sample_rtt:', sample_rtt)
412         # print('TO: %f s'%(self.timeout_interval))
413         self.set_timeout(self.timeout_interval)
```

Figure 6

The timeout is constantly recalculated based on the round-trip-time (RTT) obtained by subtracting now from the timestamp of the last received packet. An estimated RTT is then

calculated using exponential averaging, and the timeout is calculated by adding four standard deviations to the RTT.

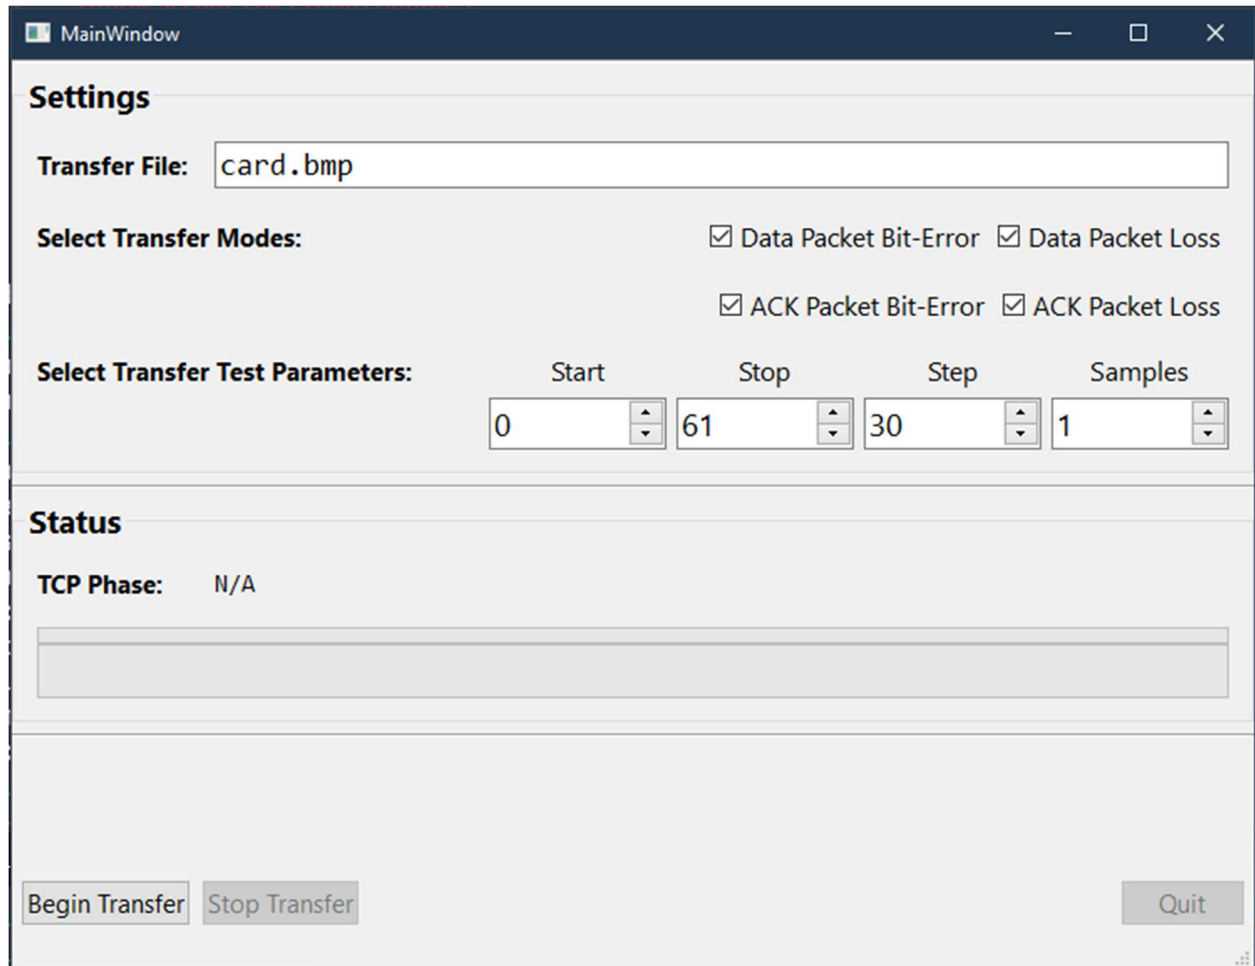
Internet Checksums

```
567 #Define bit16sum() function: Computes the packet checksum based on the checksum algorithm learned in class.
568 def bit16sum(self, data)->bytearray:
569     checksum = 0 #Initialize checksum variable, set equal to zero
570     for i, byte in enumerate(data): #For byte in data
571         if (i % 2 == 0): #If i modulus 2 == 0, then shift data by eight bits
572             bit16 = data[i] << 8
573             if (i+1 < len(data)): #If the next byte is less than the length of the data, then compute addition w/
                carryover
                bit16 = (data[i] << 8) | (data[i+1])
574             checksum = (checksum + bit16) & 0xFFFF #Compute 1s complement (addition result AND'd with 0xFFFF mask) to
                yield final checksum
575         return checksum.to_bytes(2, 'big')
```

Figure 7

A standard 16-bit checksum is calculated.

Graphical User Interface



The GUI allows the user to begin a file transfer and specify the parameters of the transfer. Two progress bars are present, one shows how far along the current file is while the second one shows how complete the entire process is.

Charts and Analysis

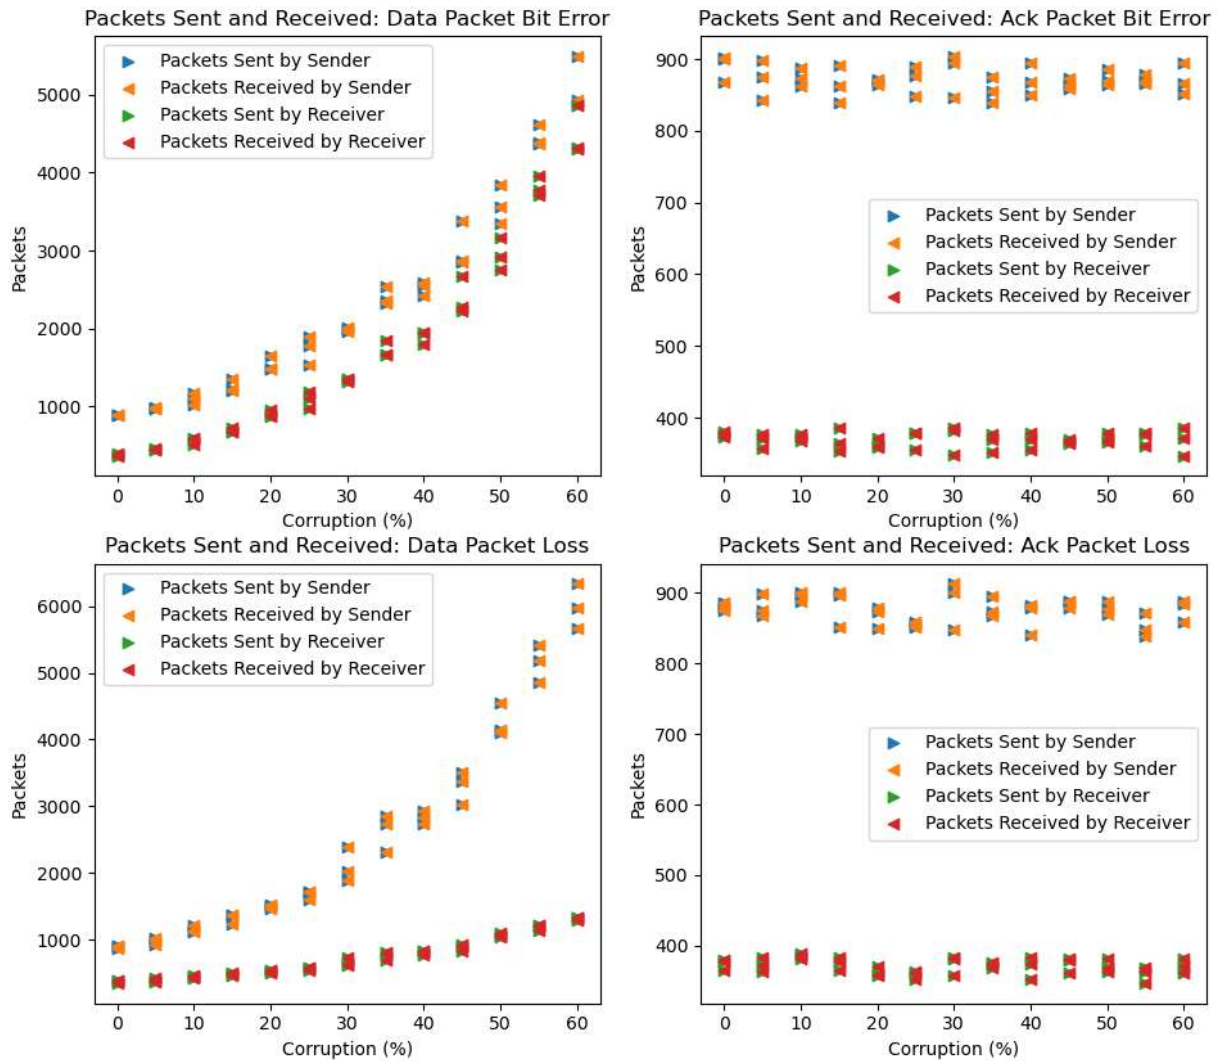


Figure 8

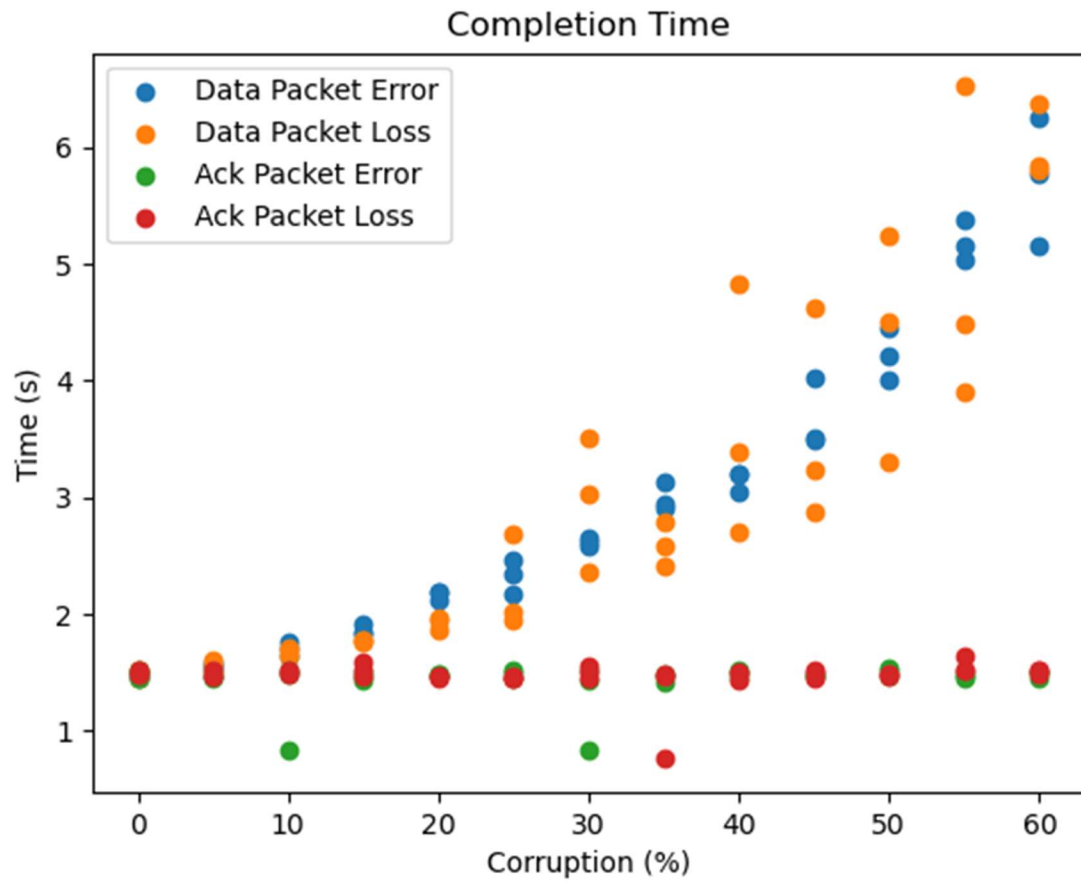


Figure 9




Figure 10




Figure 11




Figure 12




Figure 13

Figure 8 shows how many packets are sent and received for each type of corruption. Note that the number of sent packets includes dropped packets, since they appear to be sent from the sender's perspective.

Figure 9 Completion Time shows the time for completion of each round of file transfer. Three samples have been performed for each corruption-corruption type combination. The protocol is mostly impervious to acknowledgement losses since they are cumulative, and extras end up being sent anyway. Data packet errors and losses cause similar amounts of delay, indicating that the calculated timeouts are likely ideal. The high variability in transfer time during data packet loss corruption is notable.

Figures 10, 11, and 12 show the timeout over time. Recording of the timeout value begins when it is uninitialized (0), allowing us to see the initial estimate of around 1 second. It then finds equilibrium at around one-hundredth of one second. Figure 10 shows a typical connection, while Figure 11 shows a 60% corrupted line. Figure 12 offers a zoomed-in view of Figure 11 to see in detail how the timeout reacts to changes in RTT.

Figure 13 shows the window size over time for one file transfer. Note that the window size data for this graph is sampled more than once per window size adjustment, resulting in consecutively equivalent values.

References

- [1] J. Kurose, “3.4.1 Building a Reliable Data Transfer Protocol,” in *Computing Networking: A Top-Down Approach*, 7th ed., K. Ross, Ed. Pearson.
- [2] “Qt for Python (PyQt5),” *Qt for Python*. [Online]. Available: <https://doc.qt.io/qtforpython/>. [Accessed: 04-Dec-2022].
- [3] “Socket - low-level networking interface,” *Python 3.10.7 documentation*. [Online]. Available: <https://docs.python.org/3/library/socket.html>. [Accessed: 29-Sep-2022].
- [4] “Random - generate pseudo-random numbers,” *random - Generate pseudo-random numbers - Python 3.11.0 documentation*. [Online]. Available: <https://docs.python.org/3/library/random.html>. [Accessed: 30-Oct-2022].
- [5] “Time - Time Access and conversions,” *time - Time access and conversions - Python 3.11.0 documentation*. [Online]. Available: <https://docs.python.org/3/library/time.html>. [Accessed: 30-Oct-2022].
- [6] “Collections - Container Datatypes,” *collections - Container datatypes - Python 3.11.0 documentation*. [Online]. Available: <https://docs.python.org/3/library/collections.html>. [Accessed: 04-Dec-2022].
- [7] “Sys - system-specific parameters and functions,” *sys - System-specific parameters and functions - Python 3.10.7 documentation*. [Online]. Available: <https://docs.python.org/3/library/sys.html>. [Accessed: 29-Sep-2022].
- [8] “Os.path - common pathname manipulations,” *os.path - Common pathname manipulations - Python 3.11.0 documentation*. [Online]. Available: <https://docs.python.org/3/library/os.path.html>. [Accessed: 30-Oct-2022].
- [9] “Matplotlib API Reference,” *API Reference - Matplotlib 3.6.0 documentation*. [Online]. Available: <https://matplotlib.org/stable/api/index.html>. [Accessed: 30-Oct-2022].
- [10] “Signal - set handlers for asynchronous events,” *signal - Set handlers for asynchronous events - Python 3.10.7 documentation*. [Online]. Available: <https://docs.python.org/3/library/signal.html>. [Accessed: 29-Sep-2022].